

University of Padova

DEPARTMENT OF MATHEMATICS "TULLIO LEVI-CIVITA"

BACHELOR'S DEGREE IN COMPUTER SCIENCE



Explainable Machine Learning through Large Language Models: Analysis and Prompt Design

Bachelor's Thesis

Supervisor

Prof. Michele Scquizzato

Undergraduate

Lorenzo Soligo

ID number 2101057

TO DO

— TO DO

Dedicato a ...

Abstract

As **Machine Learning (ML)** models become increasingly prevalent in business applications, ensuring their transparency and reliability through **Explainable Artificial Intelligence (XAI)** is critical. This project, conducted in collaboration with Zucchetti S.p.A., investigates the interpretability and comprehensibility of various machine learning algorithms.

The research systematically evaluates a spectrum of predictive models, ranging from fundamental **Regression** and **Classification** models—such as Linear Regression, Logistic Regression, Support Vector Machines, and Decision Trees—to complex ensemble methods, including Random Forest, XGBoost and Symbolic Regression. The analysis explores both intrinsic model transparency and post-hoc explainability techniques, utilizing metrics like **SHapley Additive exPlanations (SHAP)** to interpret predictions and feature importance. Furthermore, the methodology encompasses the examination of mathematical foundations, internal knowledge representations, and the impact of ensemble techniques on both model performance and explainability. Real-world validation is performed using datasets such as Student Salary Prediction.

A primary objective of this work is bridging the gap between technical accuracy and human-understandable explanations for non-expert stakeholders. To achieve this, the project leverages advanced Prompt Engineering principles. Tailored prompts for **Large Language Model (LLM)** are developed for each analyzed algorithm. These prompts are explicitly designed to automatically generate human-readable explanations of algorithmic behaviors and predictions. Ultimately, the project delivers production-ready implementations and **LLM** integration adapters , demonstrating how the synthesis of traditional **XAI** methodologies with modern language models can significantly enhance the transparency, trust, and responsible deployment of AI systems.

TO DO

— *TO DO*

Acknowledgements

TO DO Prof. Michele Squizzato .

TO DO

TO DO

Padova, June 2026

Lorenzo Soligo

Contents

1	Introduction	1
1.1	Company	1
1.2	Stage idea	2
1.3	Project goals	2
1.4	Thesis structure	3
2	Background knowledge	4
2.1	Algorithmic analysis structure	4
2.2	Explainability	5
2.2.1	Explainability dimensions	5
2.2.2	Explanation properties and factors that facilitate understanding	5
2.2.3	SHAP analysis	6
2.3	Prompt Engineering and LLMs	7
2.3.1	Expert Personas	7
2.3.2	Familiar formats	7
2.3.3	Chain-of-Thought	9
2.3.4	Zero-Shot prompting	9
2.3.5	Multimodal prompt	10
3	ML algorithms analysis	12
3.1	Linear regression	12
3.1.1	Mathematical model	12
3.1.2	Time complexity	12
3.1.3	Spacial complexity	13
3.1.4	Internal representation	13
3.1.5	Data assumptions	14
3.1.5.1	Preprocessing	14
3.1.6	Predictive performance and limitations	14
3.1.7	Metrics for prediction quality	15
3.1.7.1	R^2 (Determination Coefficient)	15
3.1.7.2	$\bar{R}^2$ (Adjusted R^2)	15
3.1.7.3	Feature Importance (t-statistic)	15
3.1.7.4	p-value	16

3.1.7.5	Mallows' Cp	16
3.1.7.6	Root Mean Squared Error (RMSE)	16
3.1.7.7	Mean Absolute Error (MAE)	16
3.1.8	Diagnostic plots	16
3.1.8.1	Actual vs Predicted	16
3.1.8.2	Histogram of residuals distribution	17
3.1.8.3	Q-Q Plot (Quantile-Quantile)	17
3.1.8.4	Residuals vs Fitted Values	17
3.1.9	Explainability and interpretability metrics	18
3.1.9.1	Feature Effect	18
3.1.9.2	Weight Plot	18
3.1.10	Explainability limitations	18
3.2	Logistic regression	19
3.2.1	Mathematical model	19
3.2.2	Time complexity	19
3.2.3	Spacial complexity	20
3.2.4	Internal representation	20
3.2.5	Data assumptions	20
3.2.5.1	Preprocessing	21
3.2.6	Predictive performance and limitations	21
3.2.7	Metrics for prediction quality	22
3.2.7.1	Confusion Matrix	22
3.2.7.2	Accuracy	22
3.2.7.3	Sensitivity (Recall / True Positive Rate)	22
3.2.7.4	Specificity (True Negative Rate)	22
3.2.7.5	Precision	23
3.2.7.6	F1-Score	23
3.2.7.7	ROC Curve and AUC	23
3.2.7.8	Z-Statistic e p-value	24
3.2.7.9	Diagnostic plots	24
3.2.8	Feature Effect	24
3.2.9	Weight Plot	24
3.2.10	Odds Ratio	24
3.2.11	Explainability limitations	25
3.3	Support Vector Machine (SVM)	25
3.3.1	Mathematical model	25
3.3.2	Hard-Margin SVM (Linearly Separable Data)	25
3.3.3	Soft-Margin SVM (Non Linearly Separable Data)	26

3.3.4	Kernel SVM (Trasformazione Non Lineare)	26
3.3.4.1	Kernel Trick	27
3.3.4.2	Kernel Comuni	27
3.3.4.3	Rappresentazione della Soluzione	27
3.4	Complessità Computazionale	28
3.4.1	Training	28
3.4.1.1	Hard-Margin SVM (Lineare)	28
3.4.1.2	Soft-Margin SVM (con variabili di slack)	28
3.4.1.3	Kernel SVM	28
3.4.1.4	Ottimizzazioni Pratiche	28
3.4.2	Inference (Previsione)	29
3.4.3	Memoria	29
3.4.4	Scalabilità: Conclusioni	29
3.5	Rappresentazione Interna	29
3.5.1	Struttura del Modello	29
3.5.2	Implicazioni per la Spiegabilità	29
3.6	Vincoli sui Dati	30
3.6.1	Linearità (Hard-Margin SVM)	30
3.6.2	Linearità (Soft-Margin SVM)	30
3.6.3	Scalabilità	30
3.6.4	Bilanciamento delle Classi	30
3.6.5	Normalizzazione delle Feature	31
3.6.6	Nessun'altra Assunzione Strutturale	31
3.7	Capacità Predittive	31
3.7.1	Punti di Forza	31
3.7.2	Punti di Debolezza	31
3.8	Metriche per la Confidenza	32
3.8.1	Metriche di Classificazione Standard	32
3.8.2	Distanza dall'Iperpiano	32
3.8.3	Platt Scaling (Probabilità Posteriori)	33
3.9	Metriche per la Comprensione e Spiegabilità	33
3.9.1	1. Identificazione dei Support Vectors	33
3.9.2	2. Pesi dei Support Vectors (α_i)	33
3.9.3	3. Feature Importance (via Permutation o SHAP)	34
3.9.4	4. Visualizzazione della Separazione (2D/3D)	34
3.9.5	5. Analisi Locale: Contrastive Explanation	34
3.10	Limiti di Predizione	34
3.10.1	Non Probabilistico per Default	34

3.10.2	Sensibilità al Tuning di Iperparametri	35
3.10.3	Sensibilità alla Normalizzazione delle Feature	35
3.10.4	Performance su Dataset Sbilanciate	35
3.10.5	Inefficienza su Dataset Grandi	35
3.11	Limiti di Spiegabilità	35
3.11.1	Opacità della Trasformazione Non Lineare (Kernel)	35
3.11.2	Interpretazione Limitata di α_i	35
3.11.3	Dipendenza da Support Vectors Arbitrari	36
3.11.4	Scarsa Tracciabilità Globale	36
3.11.5	Nessuna Spiegazione Naturale per Feature Importanza	36
3.12	Confronto con altri Algoritmi	37
3.12.1	Vs. Logistic Regression	37
3.12.2	Vs. Decision Tree	37
3.12.3	Vs. Neural Networks	38
3.12.4	Vs. Random Forest / Gradient Boosting	38
3.13	Varianti e Estensioni	39
3.13.1	1. Multi-Class SVM	39
3.13.2	2. Support Vector Regression (SVR)	39
3.13.3	3. Uno-Classe SVM (One-Class SVM)	39
3.13.4	4. Relevance Vector Machine (RVM)	39
3.14	Raccomandazioni Pratiche	39
3.14.1	Quando Usare SVM	39
3.14.2	Tuning Consigliate	40
3.15	Prompt	40
3.16		40

4 Design and code

implementation	42
4.1 Requirements	42
4.2 Technologies and tools	43
4.2.1 Python	43
4.2.2 Markdown	43
4.2.3 Pandas, Matplotlib, Seaborn, Numpy, SHAP	43
4.2.4 Scikit-learn	43
4.2.5 Optuna	44
4.2.6 Streamlit	44
4.2.7 GitHub	44
4.2.8 Typst	44

4.3	Design	44
4.3.1	Guiding principles	45
4.3.2	Applied design patterns	45
4.3.2.1	Strategy	45
4.3.2.2	Template method	45
4.3.2.3	Facade	46
4.3.2.4	Registry	46
4.4	Code implementation	46
5	Prompt Engineering	48
6	Conclusion	50
6.1	Consuntivo finale	50
6.2	Raggiungimento degli obiettivi	50
6.3	Conoscenze acquisite	50
6.4	Valutazione personale	50
	Bibliography	52
7	Glossary	54

List of Figures

1.1	Zucchetti S.p.A. logo.	1
2.1	SHAP library logo.	7
2.2	Formats for prompting. Source: @formats-llms.	9
3.1	Histogram of residuals distribution example for linear regression.	17
3.2	Q-Q Plot of residuals example for linear regression.	17
3.3	Weight Plot of coefficients example for linear regression.	18
3.4	Confusion matrix of a logistic regression model.	22
3.5	ROC curve of a logistic regression model.	23

List of Tables

1.1	Table of the project timeline.	2
1.2	Table of project goals.	3
3.1		37
3.2		37
3.3		38
3.4		38
4.1	Table of requirements for the analysis pipeline.	42

Chapter 1

Introduction

Introduction to the stage idea, goals, company and thesis organization. This chapter is meant to provide the necessary context for the reader to understand the motivation and the structure of the thesis.s

1.1 Company

Zucchetti S.p.A. is a leading Italian software company specializing in providing comprehensive IT solutions for businesses. Founded in 1978, Zucchetti has grown to become one of the largest software providers in Italy, offering a wide range of products and services that cater to various industries, including manufacturing, retail, healthcare, and public administration.

With a strong focus on innovation and customer satisfaction, Zucchetti is committed to helping organizations optimize their operations and achieve their business goals through cutting-edge technology. Exactly from this commit, the company has been actively involved in the development and implementation of [ML](#) and [Artificial Intelligence \(AI\)](#) solutions, aiming to enhance the capabilities of their software offerings and provide more value to their clients.



Zucchetti S.p.A. logo.

1.2 Stage idea

The internship project is centered around the exploration of XAI techniques, with a specific focus on the interpretability and comprehensibility of machine learning algorithms. The primary objective is to analyze a variety of predictive models, ranging from basic regression and classification algorithms to more complex ensemble methods, in order to evaluate their transparency and explainability. The idea is to increase the understanding of the results of the models, and to make them more accessible to non-expert stakeholders, by leveraging advanced Prompt Engineering (PE) principles to develop tailored prompts for LLMs that can automatically generate human-readable explanations of algorithmic behaviors and predictions.

The timeline of the project is structured as follows:

Week	Task
1st	Analysis of Linear Regression, Logistic Regression, Support Vector Machines.
2nd	Analysis of Decision Trees, Random Forests, XGBoost.
3rd	Code implementation and Prompt Engineering for the analyzed algorithms.
4th	Evaluation and documentation of the results.
5th	Deep dive in Random Forests and research of real-world dataset.
6th	Analysis of Symbolic Regression.
7th	Code implementation and Prompt Engineering for the analyzed algorithms.
8th	Evaluation and documentation of the results, final report writing.

Table of the project timeline.

1.3 Project goals

The goals follow a notation to distinguish between necessary (N), desirable (D) and optional (O) goals. In the planning phase the goals were defined as follows:

Goal	Description
O-01	Complete and profound comprehension of the choosen algorithms, their mathematical foundations and their internal knowledge representation.
O-02	Comprension of interpretability and explainability techniques in machine learning.
O-03	Comprehension of algorithms design techniques.
O-04	Prompt generation for Large Language Models to generate human-readable explanations.
O-05	Comprehension and application of Prompt Engineering principles in the context of XAI.
D-01	Creation of a metric to evaluate the explainability of the analyzed algorithms.
D-02	Automatization of the analysis pipeline from dataset to LLM requests.
F-01	Application in real-world scenarios of the interpretability and explainability of Machine Learning models.
F-02	Explainability study of semi-supervised and unsupervised algorithms.

Table of project goals.

1.4 Thesis structure

The second chapter describes the basic concepts that fuel the project. The concepts spread from the mathematical foundations of the analyzed algorithms, the basics of explainability, to the principles of Prompt Engineering and LLMs.

The third chapter analyzes the ML algorithms, focusing on their capability, limitation and interpretability.

The fourth chapter describes the design and implementation of the analysis pipeline from dataset to LLM requests.

The fifth chapter discusses prompt engineering techniques and their application in the context of XAI.

The sixth chapter presents the conclusions of the study, the findings and the goals assessment.

Chapter 2

Background knowledge

In this chapter, we will provide the necessary background knowledge to understand the project. We will cover the mathematical foundations of the analyzed algorithms, the basics of explainability, and the principles of Prompt Engineering and LLMs.

2.1 Algorithmic analysis structure

The analysis of the algorithms is structured in a systematic way to ensure a comprehensive evaluation of their interpretability and comprehensibility as well as their functioning. The structure includes the following components:

1. **Mathematical foundations:** detailed examination of the mathematical principles underlying each algorithm, including their assumptions and theoretical properties. The focus is on the understanding of the specific **Objective Function** that the algorithms optimize, the regularization techniques they use, and the mathematical properties that govern their behavior.
2. **Computational complexity:** analysis of the computational requirements of each algorithm, including time complexities and scalability regarding both training and inference.
3. **Internal representation:** exploration of how each algorithm internally represents data and makes decisions, especially regarding **Categorical Features**.
4. **Data assumptions:** investigation into the assumptions each algorithm makes about the data, such as linearity, independence of features, and distributional assumptions.
5. **Predictive performance and limitations:** evaluation of the predictive performance of each algorithm on relevant datasets, considering the data assumptions and mathematical foundation of the model.
6. **Metrics for prediction quality:** presentation of the most relevant metrics for evaluating the models performance, based on the task.
7. **Metrics for interpretability:** presentation of the most relevant metrics for evaluating the models interpretability, based on the task.
8. **Explainability limitations:** discussion of the limitations of explainability techniques for each algorithm, including potential pitfalls and challenges in interpreting their predictions.

2.2 Explainability

Explainability (or interpretability) of a ML model is the ability to communicate **how and why** the model made a prediction.

Contrary to intuition, explainability it's **not a binary attribute**, but a spectrum with multiple dimensions.

2.2.1 Explainability dimensions

1. Intrinsic Transparency (Model-Level Interpretability)

The model is inherently transparent and interpretable by design. The structure directly communicates how it works, such as regression coefficients or decision tree nodes. Each prediction is fully traceable, and parameters have tangible meanings. Explanations reflect human reasoning, making them coherent from a domain perspective.

It represents the best case scenario for explainability but it is often in trade-off with performance, as more complex models tend to be less transparent but more powerful. It is generally organized in tiers: High, Medium and Low transparency.

2. Global Interpretability (Model-Level Explainability)

The model is comprehensible in its entirety, allowing to understand its overall behaviour and decision-making process. Such a level of understanding it's hardly achievable and it is usually indirectly reached by comprehending the singular components of the models instead of the model as a whole.

3. Modular Interpretability

When global interpretability is not achievable, we can still understand the model by analyzing its components separately. Exploiting the model transparency, we can understand the role of each component in the decision-making process.

This concept is clear in the case of linear methods, where it is easy to understand the contribution of the single weights or in tree based models, where the splits and the structure of the tree can be analyzed separately.

4. Local Interpretability (Instance-Level Explainability)

The model might be opaque and too complex to understand globally. However, for a specific prediction, we can generate an explanation that is locally interpretable. SHAP could be used in this context to explain a particular instance obtained a certain prediction.

2.2.2 Explanation properties and factors that facilitate understanding

For human explanations to be effective, they should have certain properties that facilitate understanding and usability. The following properties can be identified from the work of Robnik-Sikonja and Bohanec 2018 and Molnar [1]:

- **Contrastive explanations**

Comparative explanations that highlight the differences between instances can be more informative than absolute explanations. The instances confrontation makes the explanation more impactful, providing a direct idea of how different factors contribute to different outcomes.

- **Cause selection**

Not all features are equally important; explanations should focus on the most influential factors rather than listing all contributing features. A small sample of the most relevant features provide a more clear explanation than a long list of all features preventing information overload.

- **Social and contextual factors**

The social environment and the audience’s expertise level should be considered when generating explanations to ensure they are appropriate and understandable.

- **Anomalies or abnormal predictions**

Highlighting anomalies or abnormal predictions can help users understand when the model is making an unusual decision, which can be crucial for trust and debugging. Features that are not largely relevant for the majority of the predictions but that have a large impact on specific predictions should be highlighted for their capability.

- **Fidelity**

The explanation should accurately reflect the model’s behavior and not introduce misleading information. A high-fidelity explanation ensures that users can trust the insights provided and make informed decisions based on them, provided that the model is accurate in its predictions.

- **General and probable**

In the absence of anomalies of specific cases, a good explanation should rely on general and probable causes that are likely to be true for most instances. This parameter can be quantified by the feature support:

$$\text{Feature Support}_i = \frac{\#\text{training instances with similar values to } x_i}{\#\text{total training instances}}$$

Generality can be misleading, as it may lead to explanations that are too broad and not specific enough to be useful. It’s important to balance generality with the impact of the explanation, ensuring that it provides meaningful insights without being overly vague. This is why usually **abnormal features** are highlighted, as they can provide more specific and impactful explanations.

2.2.3 SHAP analysis

SHAP (SHapley Additive exPlanations) is a method for explaining the output of machine learning models by attributing the prediction to each feature. It is based on the concept of Shapley values from cooperative game theory, which provide a way to fairly distribute the “payout” (in this case, the prediction) among the “players” (the features) based on their contribution to the prediction. SHAP values can be used to understand the importance of each feature for a specific

prediction (local interpretability) or for the model as a whole (global interpretability). SHAP provides a unified framework for interpreting various types of models and can be applied to both linear and non-linear models, making it a powerful tool for explainable AI.



SHAP library logo.

2.3 Prompt Engineering and LLMs

The following sections will present the principles used in the design of the prompts for the LLMs to generate human-readable explanations of the analyzed algorithms.

2.3.1 Expert Personas

The expert personas techniques consists in the design of prompts that describes the role that the LLM should play in the task. This technique first should guide the model in generating a more precise, coherent and relevant output.

Though the expert personas seems to be linked to a loss of accuracy in multiple benchmarks in the recent studies by Zizhao Hu [2], the principal use of the LLMs in this project is results evaluation and text generation. The data and metrics do not have to be generated and have only to be read directly from the prompt. The paper specifically talks about “*For tasks that depend on pretrained knowledge retrieval accuracy*” [2]. It’s reasonable to think that in such a context, loss of accuracy is not a critical issue. In the same paper, an alignment of behavior and style to the described personas is observed. This drift in behavior is critical in the task of describing the results of an algorithm in a human-readable way, as the style of the explanation is fundamental to make it accessible to non-expert stakeholders.

2.3.2 Familiar formats

The extensive use of LLMs have led to the emergence of some familiar formats that are widely used in the design of prompts.


The principal features that this formats present, and that makes them particularly effective, are the following:

- Widely use in training: the more a format is used in the training of the model, the more likely is to be effective in the interpretation of content in that format.
- Clear structure: the format should have a clear and recognizable structure that allows the model to easily identify the different components of the prompt and their relationships.
- Token efficiency: the format should be designed in a way that minimizes the number of tokens used, while still conveying all the necessary information. This is particularly important in the context of LLMs, as the number of tokens used can have a significant impact on the performance and cost of the model.

From the multitude of formats in existence, some emerged for their particular accuracy and efficiency. The most valueable one is **Markdown**, which has a clear and recognizable structure that allows the model to easily identify the different components of the prompt and their relationships. The use of markdown also allows to create a clear and organized structure for the generated explanations, making them more readable and accessible to non-expert stakeholders. Another accurate format is YAML, which offers a more readable formatting than JSON and XML while still being well structured.

An alternative like CSV can be used, as it is a widely used format for tabular data and is likely to be well understood by the model. However, it's important to ensure that the CSV format is properly structured and that the necessary information is included in the prompt to allow the model to generate accurate and relevant explanations.

What follows is a summarizing image where the 11 formats taken in exam in the study[3] are presented. The y-axis is a scale of Accuracy while the x-axis a token consumption scale. It's clear that in this particular case of study, only OpenAI's GPT-4.1 nano has been tested, Markdown emerges as the most accurate format, while maintaining good token usage.



Formats for prompting. Source: @formats-llms.

2.3.3 Chain-of-Thought

Chain of Thought (CoT) prompting is a technique that consists in the design of prompts that guide the LLMs to generate a step-by-step reasoning process to reach the final answer. This technique is particularly useful in the context of this project, as it allows to generate explanations that are more coherent and relevant to the analyzed algorithms.

The additional value of the guided, step-by-step reasoning has been proven in multiple benchmarks[4], and it is particularly useful in the context of this project, as it allows to generate explanations that are more coherent and relevant to the analyzed algorithms.

In the particular context of XAI and non-technical audiences because it enhances the model's ability to provide detailed and understandable explanations. The generated answer is not just a final result, but a comprehensive explanation that breaks down the reasoning process into clear and logical steps.

2.3.4 Zero-Shot prompting

This technique consists in the design of prompts that do not include any example of the expected output. It's the simplest kind of prompting because it erases the problem of having to provide training examples. Such method is particularly necessary in the context due to the fact that an example of data and metrics for reference could be too specific and not generalized enough. An additional set of data could also confuse the model and make it less accurate or biased on

the actual result of the analysis.

The zero-shot prompting also revealed to be quite effective with some additional techniques such as the use of **Chain of Thought** (CoT) prompting [5].

2.3.5 Multimodal prompt

The multimodal prompting technique consists in the design of prompts that include multiple modalities, such as text, images, tables, etc. This technique is particularly useful in the context of this project, as it allows to provide the model with a richer and more comprehensive context for generating explanations.

The use of multimodal prompts can enhance the model’s ability to comprehend and analyze the results of the algorithms, as it can leverage the information provided in different formats to generate more accurate and relevant explanations. In fact, *“visual prompts effectively interact with the textual prompts, enhancing the alignment between modalities and thereby improving the model’s performance on the zero-shot instruction learning”* [6].

Chapter 3

ML algorithms analysis

In this chapter, we will analyze the machine learning algorithms used in the project, looking at the mathematical foundations, the data requirements, the predictive performance, and the interpretability of the results.

This initial analysis will guide the design and implementation of the pipeline as well as the creation of the specific prompts for the [LLMs](#) to generate human-readable explanations of the results.

3.1 Linear regression

3.1.1 Mathematical model

The linear regression model is the simplest form of regression. It assumes a linear relationship between the input features and target variable. The mathematical representation of the linear regression model is:

$$y = \sum_{i=0} \beta_i x_i \forall i \in F$$

Where F is the set of features, β are the calculated weights for each of the features, and y is the target variable.

The goal of the linear regression is to find the optimal weights β that minimize the error between the predicted values and the actual target variable. This is typically done by minimizing the ordinary least squares loss function between the predicted values and the actual target variable. The optimization problem can be formulated as:

$$\hat{\beta} = \operatorname{argmin}_{\beta_0 \dots \beta_p} \sum_{i=1}^n (y^{(i)} - (\beta_0 + \sum_{j=1}^p \beta_j x_j))^2$$

3.1.2 Time complexity

Mainly there are two approaches to solve this optimization problem: the analytical solution and the iterative optimization methods (e.g. [Gradient Descent \(GD\)](#)). The analytical solution is given by the normal equation:

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

In this formulation, X is an $n \times p$ matrix where n is the number of instances and p is the number of features. The matrix X is augmented with a column of 1s to account for the intercept term β_0 . The [computational complexity](#) of the analytical solution is dominated by

the matrix multiplication and inversion operations. In fact, the cost for the multiplication of matrixes is $O(p^2n)$ and the cost for the inversion of a $p \times p$ matrix is $O(p^3)$. Therefore, the overall **computational complexity** of the analytical solution is $O(p^2n + p^3) = O(p^3)$, which for vast datasets, over 10000 records, becomes prohibitive.

The iterative methods, such as **GD**, calculate the gradient of the loss function with respect to the weights and update the weights iteratively until convergence. The **computational complexity** of the gradient calculation is $O(pn)$ per iteration, and the number of iterations required for convergence can vary depending on the learning rate and the specific dataset. However, in practice, iterative methods can be more efficient than the analytical solution for large datasets, as they do not require matrix inversion and can converge faster with appropriate hyperparameter tuning. Other methods like **Stochastic Gradient Descent (SGD)** can further reduce the computational cost by approximating the gradient using a subset of the data at each iteration. On the other side, for smaller dataset, analytical solution offers convergence in a single step, without the need for hyperparameter tuning, and guarantees a global optimal solution.

For **inference**, the **computational complexity** is $O(p)$ per instance, as it involves a simple dot product between the feature vector and the weight vector.

3.1.3 Spacial complexity

The spacial complexity of the linear regression model is determined by the need to store the input data, the weight vector, and any intermediate matrices used in the analytical solution. Specifically, the analytical solution requires storing the $n \times p$ matrix X , the $p \times p$ matrix $X^T X$, and the p -dimensional weight vector β . Therefore, the overall spacial complexity of the linear regression model is dominated by the storage of the input data and the intermediate matrices, resulting in a spacial complexity of $O(np + p^2)$.

For the **GD** method, the spacial complexity is reduced to $O(np + p)$, as it does not require storing the intermediate matrix $X^T X$.

For **inference**, the **computational complexity** is $O(p)$ per instance, as it only stores the weight vector and the feature vector.

3.1.4 Internal representation

Linear regression represents the resulting weight of the features as a vector of weights, $\beta = [\beta_0, \beta_1, \dots, \beta_p]$.

Specific encoding is needed for **categorical features**, which are typically handled through one-hot encoding, where each category is represented as a binary feature. This can lead to an increase in the number of features and potential multicollinearity issues if not handled properly.

In regard of **explainability**, the internal representation of linear regression is straightforward and transparent, which makes it one of the most interpretable machine learning models. The weights directly indicate the strength and direction of the relationship between each feature

and the target variable. This allows for a clear understanding of how each feature contributes to the prediction, making it easier to communicate insights to stakeholders and identify important predictors in the data.

3.1.5 Data assumptions

Linear regression relies on several key assumptions about the data to ensure the validity of the model and the reliability of its predictions. These assumptions include as described by Molnar [1]:

1. **Linear constraints:** relationships between features and target must be linear. Other kind of relationships must be manually introduced and cannot be revealed automatically
2. **Residuals normality:** the residuals $\epsilon_i = y_i - \hat{y}_i$ must follow a normal distribution. Heavy violations compromise the inference.
3. **Homoscedasticity:** residuals must have constant variance across all levels of the features. In practice, this assumption is **frequently violated**. Example: in real estate, the price of very large houses is extremely variable, while that of small houses is concentrated.
4. **Indipendence of measurements:** instances don't have to be correlated. Dipendent datas, such as time series, violate this assumption and as such should not be investigated with linear regression.
5. **Feature fisse:** le feature devono essere considerate come costanti, non soggette a errori di misurazione significativi. Se le feature contengono errore di misura, i coefficienti sono distorti (attenuation bias)
6. **Absence of multicollinearity:** Feature should not be highly correlated. Multicollinearity causes numerical instability during the inversion of $X^T X$ e weights inflation in absolute value. Using **GD** the geometrical properties of the loss function changes and leads to very slim vallies causing an important reduction of the learning rate.

3.1.5.1 Preprocessing

The data assumption lead to specific preprocessing steps to determine the suitability of the data for linear regression and to improve the performance of the model. These steps include:

1. **Multicollinearity identification:** calculation of the **Correlation Matrix** between features using the Pearson's coefficient or VIF (Variance Inflation Factor) to identify highly correlated features. $VIF_j = \frac{1}{1-R_j^2}$

Where R_j^2 is the coefficient of determination for the regression of feature j against all other features. More on R_j^2 [here](#).

3.1.6 Predictive performance and limitations

As discussed, the linear regression imposes several constraints on the data and model performance. Very good performances are achieved whenever linear relationships exist between

features and target both in accuracy and efficiency.

However, in real-world scenarios, these conditions are often violated, leading to poor predictive performance. The model is particularly sensitive to **outliers**, which can disproportionately influence the weights and lead to skewed predictions. Additionally, linear regression is **not suitable for capturing complex, non-linear relationships** in the data, which limits its applicability in many real-world problems where such relationships are common. Finally, the model's performance can degrade significantly when the number of features is large relative to the number of observations, leading to overfitting and poor generalization to new data. This limitation is intensified by the presence of **categorical features**.

3.1.7 Metrics for prediction quality

What follows is a list of most relevant metrics for evaluating the predictive performance of linear regression models, based on the task and the data assumptions.

3.1.7.1 R^2 (Determination Coefficient)

R^2 quantifies how much the model explains the total variance of the data. It ranges from 0 to 1, where 0 is a model that cannot explain the data and 1 is a perfect adherence.

$$R^2 = 1 - \frac{SSE}{SST}$$

Where:

- $SSE = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$ is the sum of squared residuals, quantifying the variance unexplained by the model
- $SST = \sum_{i=1}^n (y^{(i)} - \bar{y})^2$ is the total variance

3.1.7.2 $\bar{R}^2$ (Adjusted R^2)

R^2 tends to increase with the number of features, even if those features are not useful. Adjusted R^2 penalizes the addition of non-informative features.

$$\bar{R}^2 = 1 - (1 - R^2) \frac{n - 1}{n - p - 1}$$

Where n is the number of instances and p is the number of features.

3.1.7.3 Feature Importance (t-statistic)

$t_{\hat{\beta}_j}$ measures the statistical significance of each coefficient, calculated as the weight scaled by its standard error:

$$t_{\hat{\beta}_j} = \frac{\hat{\beta}_j}{SE(\hat{\beta}_j)}$$

Intuitively, the higher the absolute value of a coefficient is, the more the feature is statistically significant in predicting the target variable.

As intuitively, the higher the variance of the coefficient is, the less the feature is significant, as the model is more uncertain about the true value of the coefficient.

3.1.7.4 p-value

p-value is a measure of the probability of “obtaining the observed data under the null hypothesis of a statistical test”[7]. In the context of linear regression, the null hypothesis is that the true coefficient β_j is equal to zero, meaning that the feature does not have a significant impact on the target variable. The p-value for each coefficient is calculated based on the t-statistic and indicates the probability of observing such an extreme value for $t_{\hat{\beta}_j}$ if the null hypothesis were true. A common convention is to consider a p-value less than 0.05 as statistically significant, suggesting that there is strong evidence against the null hypothesis and that the feature likely has a meaningful relationship with the target variable.

3.1.7.5 Mallows' Cp

Mallows' Cp is a model selection metric that balances model fit and complexity, calculated as:

$$Cp = \frac{SSE}{\hat{\sigma}^2} - n + 2p$$

Where $\hat{\sigma}^2$ is the estimate of residuals variance of the complete model. It's used to reduce the problem of overfitting, reducing the number of features.

3.1.7.6 Root Mean Squared Error (RMSE)

RMSE measures the magnitude of the errors, in the same scale as the target variable:

$$RMSE = \sqrt{\frac{SSE}{n}}$$

3.1.7.7 Mean Absolute Error (MAE)

MAE measures the average magnitude of the errors, without considering their direction:

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

3.1.8 Diagnostic plots


The use of diagnostic plots is crucial to visually assess the assumptions of linear regression and to identify potential issues with the model fit.

3.1.8.1 Actual vs Predicted

Scatter plot with actual values y_i on the y-axis and predicted values $\hat{y}_i$ on the x-axis. $\hat{y}_i$. If the model fits well, the points should be concentrated around the diagonal line $y = \hat{y}$. Deviations from this pattern can indicate various issues with the model fit.


3.1.8.2 Histogram of residuals distribution

Distribution of the residuals $\epsilon_i = y_i - \hat{y}_i$. It's especially useful to identify violations of the normality assumption. A normal distribution of residuals is expected for valid inference, and deviations from this pattern can indicate issues with the model fit or the presence of outliers.



Histogram of residuals distribution example for linear regression.

3.1.8.3 Q-Q Plot (Quantile-Quantile)



Q-Q Plot of residuals example for linear regression.

Another way to check the normality of residuals is through a Q-Q plot, which compares the quantiles of the residuals to the quantiles of a normal distribution. If the residuals are normally distributed, the points in the Q-Q plot should approximately follow a straight line. Deviations from this line, especially at the tails, can indicate non-normality of the residuals, which may affect the validity of statistical inference based on the model.

3.1.8.4 Residuals vs Fitted Values

Scatter plot with fitted values $\hat{y}_i$ on the y-axis and residuals $\epsilon_i = y_i - \hat{y}_i$ on the x-axis. $\hat{y}_i$. If the model fits well, the points should be concentrated around the diagonal line $y = \hat{y}$. This plot can be used to identify violation in regression assumptions ([Section 3.1.5](#)). Increasing dispersion for example, show heteroschedasticity, while systematic patterns (e.g. a curve) indicate non-linearity that the model cannot capture.

3.1.9 Explainability and interpretability metrics

3.1.9.1 Feature Effect

Linear regression naturally provides a direct measure of the feature effect on the target variable through the coefficients β_j . The effect of a feature x_j on the prediction for an instance i can be calculated as:


$$\text{effect}_j^{(i)} = \beta_j x_j^{(i)}$$

The standard visualization shows a box plot of the calculated effect in the quantiles 25, 50 and 75 for each feature. This allows to quickly identify the features with the most significant effect on the predictions, as well as the distribution of the effects across the dataset.

This plot results not useful if the datas are normalized, as the effect is calculated as the product of the coefficient and the feature value, and normalization can obscure the true impact of the features on the predictions.

3.1.9.2 Weight Plot

For a direct visualization of the importance of each feature, a weight plot can be used, which shows the coefficients β_j for each feature. The coefficients indicate the strength and direction of the relationship between each feature and the target variable.



Weight Plot of coefficients example for linear regression.

3.1.10 Explainability limitations

As emerged until now, linear regression is one of the most interpretable machine learning models due to its transparent internal representation and the direct relationship between features and

predictions. The impact of each feature on the prediction can be easily understood through the coefficients, which indicate how much the prediction changes with a one-unit change in the feature, holding all other features constant.

What makes the model extremely interpretable make it limited in its predictive ability. Linearity of the relationships is as understandable as restrictive.

3.2 Logistic regression

3.2.1 Mathematical model

The logistic regression is a model mainly used for binary classification problems where the response variable is dichotomous. Similar to linear regression(Section 3.1) it uses a linear combination of the input features, but instead of directly outputting a continuous value, it applies a logistic function to map the output to a probability between 0 and 1.

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$

The output can consequently be interpreted as the probability that the input instance belongs to the positive class ($y = 1$). In particular, the probability that an instance belongs to the positive class is given by:

$$P(y^{(i)} = 1 \mid x^{(i)}) = \frac{1}{1 + \exp(-(\beta_0 + \sum_{j=1}^p \beta_j x_j^{(i)}))}$$

And the probability that it belongs to the negative class ($y = 0$) is $P(y^{(i)} = 0) = 1 - P(y^{(i)} = 1)$, following the Bernoulli distribution.

The goal of the logistic regression is to find the parameters β that best fit the data, which is done by maximizing the **likelihood** of observing the data given the parameters.

$$L(\beta; y, X) = \prod_{i=1}^n P(y^{(i)} = 1)^{y^{(i)}} \cdot (1 - P(y^{(i)} = 1))^{1-y^{(i)}}$$

For numerical stability **log-likelihood** is usually calculated:

$$\ell(\beta) = \sum_{i=1}^n [y^{(i)} \log(P(y^{(i)} = 1)) + (1 - y^{(i)}) \log(1 - P(y^{(i)} = 1))]$$

Differently from linear regression, the logistic regression does not have a closed-form solution for the parameters that maximize the likelihood so only iterative optimization algorithms like **GD** are available

3.2.2 Time complexity

As described in Section 3.1.2, the time complexity of **GD** is $O(pn)$ per iteration, and the number of iterations required for convergence can vary depending on the learning rate and the specific dataset. For **inference** the time complexity is again, $O(p)$ per instance.

Since the optimization is iterative, the overall time complexity can be less predictable than linear regression, especially if the data has features that lead to slow convergence (e.g., highly correlated features or features that cause the decision boundary to be close to the data points). However, in practice, logistic regression is often computationally efficient for moderate-sized datasets and can be scaled to larger datasets using **SGD** or mini-batch gradient descent.

3.2.3 Spacial complexity

The model stores a vector of coefficients β of size p , which requires $O(p)$ memory. During training, additional memory is needed to store the intermediate values for the optimization algorithm, which can be $O(n)$ for batch gradient descent due to the need to compute the predictions for all instances in each iteration.

For inference, the memory complexity is $O(p)$ for storing the coefficients and $O(1)$ for the input instance, leading to an overall spacial complexity of $O(p)$.

3.2.4 Internal representation

The model internally represents the values as a vector of weights, $\beta = [\beta_0, \beta_1, \dots, \beta_p]$. The presence of **categorical features** is solved as standard through one-hot encoding just like in linear regression.

For **explainability**, the model remains transparent since the coefficients still indicate the strength and direction of the relationship between each feature and the target variable, but the interpretation is less intuitive than in linear regression due to the non-linear transformation applied to the output by the logistic function. An increase in one feature in logistic regression leads to a multiplicative change in the odds of the positive class, rather than an additive change in the predicted value.

3.2.5 Data assumptions

Before discussing the data assumptions we need to clarify the role of the odds. The odds of an event is defined as the ratio of the probability of the event occurring to the probability of it not occurring:

$$\text{odds} = \frac{P(y = 1)}{P(y = 0)} = \exp\left(\beta_0 + \sum_{j=1}^p \beta_j x_j\right)$$

This formulation more easily shows how the coefficients β_j affect the predictions.

$$\frac{\text{odds}_{x_j+1}}{\text{odds}} = \exp(\beta_j)$$

Just like in linear regression, the idea of maintaining the interpretability of the model through a linear combination of the features leads to several assumptions on the data and the model performance:

1. **Linearity of the log-odds:** $\log(\text{odds}) = \beta_0 + \sum_j \beta_j x_j$. The relationship is linear in the **log-odds space**, not in the probability space. Nonlinear relationships remain uncaptured and must be added manually.
2. **Complete separation:** if a feature perfectly separates the two classes, the model cannot identify a finite weight since the algorithm will tend to push the weight to $+\infty$ or $-\infty$, diverging. Usually the solution is to erase this separating feature, since it is probably just a proxy for the target variable.
3. **Binomial distribution:** the response variable must follow a **Bernoulli distribution**.
4. **Independence of measurements:** observations should not be correlated. This means that the model is not suitable for temporal data without proper control, as well as for data with strong dependencies between observations.
5. **Fixed features:** features must be considered constants, without measurement error.
6. **Absence of multicollinearity:** highly correlated features cause coefficient instability. This is particularly relevant due to the fact that the iterative methods begin to oscillate and need very small learning rates (α).

Homoscedasticity is no more a requirement because the response can only be 0 or 1.

3.2.5.1 Preprocessing

To verify the suitability of logistic regression in the classification task for a given dataset, some methods can be used. The **correlation matrix** can be used to identify highly correlated features, which can lead to multicollinearity issues. If such features are found, they can be removed or combined to reduce redundancy and improve model stability.

For identifying anomalies in the other assumptions, plot visualisation can be used. See [Section 3.2.7.9](#) for more details.

3.2.6 Predictive performance and limitations

The imposed assumptions of logistic regression can lead to good performance whenever these assumptions are met. The probabilistic interpretation gives insight on the confidence of the prediction, which is a significant advantage in many applications. Computationally wise, it achieves fast training thanks to the iterative methods to solve the loss function.

However when the relationships are more complex than basic linear combinations, the predictions become less accurate. In context of unbalanced classes the training process can be dominated by the majority class, leading to poor performance on the minority class. Additionally, logistic regression as the linear regression is highly influenced by outliers, leading to very unstable predictions. There's then the limitation of the complete separation, which prevents the model to converge to a solution.

3.2.7 Metrics for prediction quality

The following is a list of metrics for evaluating the predictive performances of the logistic regression in the context of the classification task.


Before presenting this metrics, it is important to define some terms, used later:

- **True Positives (TP)**: instances correctly predicted as positive
- **True Negatives (TN)**: instances correctly predicted as negative
- **False Positives (FP)**: instances incorrectly predicted as positive
- **False Negatives (FN)**: instances incorrectly predicted as negative

3.2.7.1 Confusion Matrix

The confusion matrix is a table that summarizes the results of a classification task by comparing the true class labels with the predicted class labels.

The 4 different categories are **TP**, **TN**, **FP** and **FN** as described before.



Confusion matrix of a logistic regression model.

3.2.7.2 Accuracy

Accuracy measures the ratio of correct predictions to the total predictions:

$$\text{ACC} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

It is important to notice that in the context of unbalanced classes, accuracy can be misleading, as a model that always predicts the majority class can achieve high accuracy while performing poorly on the minority class.

3.2.7.3 Sensitivity (Recall / True Positive Rate)

The sensitivity, also known as recall or true positive rate, measures the ratio of correctly predicted positive instances to all actual positive instances:

$$\text{SENS} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

3.2.7.4 Specificity (True Negative Rate)

The specificity, also known as true negative rate, measures the ratio of correctly predicted negative instances to all actual negative instances:

$$\text{SPEC} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

Sensitivity and specificity are particularly important in contexts where the cost of false positives and false negatives is different, such as in medical diagnosis.

3.2.7.5 Precision

Precision measures the ratio of correctly predicted positive instances to all predicted positive instances:

$$\text{PREC} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Precision is crucial in scenarios where the cost of false positives is high, such as in spam detection or fraud detection.


3.2.7.6 F1-Score

Out of the box, the precision and recall can be in tension, as improving one often leads to a decrease in the other. The F1-score is the harmonic mean of precision and recall, providing a single metric that balances both:

$$\text{F1} = 2 \frac{\text{PREC} \cdot \text{REC}}{\text{PREC} + \text{REC}}$$

It results particularly useful in unbalanced classification problems or in situations in which both false positive and false negative are costly.

3.2.7.7 ROC Curve and AUC



ROC curve of a logistic regression model.

ROC Curve is visual representation of the True Positive Rate (Section 3.2.7.3) - False Positive Rate trade-off Positive Rate as the threshold of classification varies. The Sensitivity sits on the y-axis and False Positive Rate on the x-axis.

A model with good performance will have a curve that bows towards the top-left corner of the plot, indicating high sensitivity and low false positive rate across different thresholds. A model that predicts randomly will have a curve that follows the diagonal line.

AUC (Area Under the Curve) is exactly the area under the ROC curve, numerically quantifying the overall ability of the model to discriminate between the positive and negative classes. The AUC ranges from 0 to 1, with higher values indicating better performance and 0.5

representing random guessing.

Can be interpreted as the probability that the model will rank a randomly chosen positive instance higher than a randomly chosen negative instance.

3.2.7.8 Z-Statistic e p-value

The specular t-statistic for logistic regression is the Z-statistic, which measures the statistical significance of each coefficient in the model. It is calculated as:

$$Z = \frac{\beta_j}{\text{SE}(\beta_j)}$$

Where $\text{SE}(\beta_j)$ is the standard error of the coefficient β_j . A coefficient with a large absolute value of Z indicates that the coefficient is significantly different from zero, suggesting that the corresponding feature has a significant impact on the predicted probabilities.

The p-value associated with the Z-statistic represents the probability of observing such an extreme Z value under the null hypothesis that $\beta_j = 0$. For more details on the p-value, see [Section 3.1.7.4](#).

3.2.7.9 Diagnostic plots

Using plots to visualize the data and the model's predictions can help identify potential issues with the assumptions of logistic regression and provide insights into the model's performance.

3.2.8 Feature Effect

Similarly to Linear Regression, [Section 3.1.9.1](#), the feature effect plot rappresent the impact of a feature on the prediction. In the logistic regression case, the effect is not on the predicted value but on the predicted probability, which is more intuitive to understand for most users.

For every feature value, calculate:

$$P(y = 1 \mid x_j = v, x_{\text{other}} = \text{media})$$

Vantaggio rispetto a LR: la trasformazione sigmoide rende visibile se l'effetto è principalmente presso certi valori della feature (grafico non è una retta, è una curva).

3.2.9 Weight Plot

Likewise to linear regression, the weight plot shows the coefficients of the features. However, in logistic regression, the coefficients do not directly correspond to changes in the predicted probability, but rather to changes in the log-odds of the positive class.

3.2.10 Odds Ratio

Direct expression of the feature impact on the odds:

$$\text{OR}_j = \exp(\beta_j)$$

It represents the multiplicative change in the odds of the positive class for a one-unit increase in the feature x_j , holding all other features constant.

3.2.11 Explainability limitations

The logistic regression, while being more interpretable than many other machine learning models, has some limitations in terms of explainability. Firstly, the interpretation of the coefficients is less intuitive than in linear regression due to the non-linear transformation applied to the output by the logistic function. An increase in one feature in logistic regression leads to a multiplicative change in the odds of the positive class, rather than an additive change in the predicted value.

This multiplicative effect is even less straightforward if we consider that the effect of a feature on the predicted probability depends on the values of all the other features, making it difficult to provide a simple explanation of the effect of a single feature without considering the context of the other features.

3.3 Support Vector Machine (SVM)

3.3.1 Mathematical model

A Support Vector Machine is a classifier that finds the optimal hyperplane that separates two classes by maximizing the distance (**margin**) between the hyperplane and the closest points of each class. (ciascuna classe). In the case of bidimensional data, the hyperplane is a line; for three-dimensional data, it is a plane; for p-dimensional data, it is a hyperplane defined by:

$$\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p = 0$$

Where $\beta = [\beta_1, \dots, \beta_p]$ is the normal vector to the hyperplane, and β_0 is the intercept.

3.3.2 Hard-Margin SVM (Linearly Separable Data)

In the ideal case in which the data are perfectly and **completely separable**, the goal is to find the coefficients $\beta_0, \beta_1, \dots, \beta_p$ that maximize the **margin**. The intuition, bias, is that a hyperplane with a large margin is more **robust** to variations in the data and generalizes better to unseen data. The separation condition is then the following:

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq 1 \quad \forall i$$

Where $y_i \in \{-1, +1\}$ is the class label.

The distance between a point in the training set and the hyperplane is given by:

$$r_i = \frac{y_i(\beta_0 + \sum_{j=1}^p \beta_j x_{ij})}{\|\beta\|}$$

Where $\|\beta\| = \sqrt{\sum_{j=1}^p \beta_j^2}$ is the **euclidean norm** of the coefficient vector.

To maximize the **margin**, defined as the distance between the hyperplane and the closest points, we can express it as:

$$M = \frac{1}{\|\beta\|}$$

is equivalent to minimizing $\|\beta\|$. This formulation leads to the following optimization problem:

$$\min_{\beta, \beta_0} \frac{1}{2} \|\beta\|^2$$

Under the constraint that:

$$y_i \left(\beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) \geq 1 \quad \forall i = 1, \dots, n$$

3.3.3 Soft-Margin SVM (Non Linearly Separable Data)

In the realistic case where the data are **not linearly separable**, for the presence of noise, outliers, or overlapping classes, we allow some points to violate the margin. This can be achieved by introducing **slack variables** $\xi_i \geq 0$ that measure how much a point violates the margin:

$$y_i \left(\beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right) \geq 1 - \xi_i \quad \forall i$$

Once we allow violations, we need to penalize them in the objective function to prevent the model from simply classifying all points as the majority class. This leads to the following optimization problem:

$$\min_{\beta, \beta_0, \xi} \left[\frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^n \xi_i \right]$$

The objective function shows the importance of two components of the model. Firstly the distance of the hyperplane (first term) and secondly the violations of the margin (second term). The parameter C is a **hyperparameter of regularization** that controls the trade-off between maximizing the margin and minimizing the violations. A high value of the parameter C will prioritize minimizing the violations, rapproching the hard-margin SVM and potentially leading to a higher overfitting. A low value of C will prioritize maximizing the margin, allowing more violations.

3.3.4 Kernel SVM (Trasformazione Non Lineare)

Il limite principale di SVM lineare è che funziona solo se i dati sono (approssimativamente) **linearmente separabili**. Per dati con pattern non lineari, SVM usa il **kernel trick**.

Idea: trasformare i dati in uno spazio di dimensionalità superiore (possibilmente infinita) dove diventano linearmente separabili. Tuttavia, computare esplicitamente la trasformazione è costoso. Il **kernel trick** consente di fare questo implicitamente.

3.3.4.1 Kernel Trick

Anziché trasformare i dati esplicitamente $\phi(x_i) = [f_1(x_i), f_2(x_i), \dots, f_m(x_i)]$ e poi calcolare prodotti scalari $\phi(x_i)^T \phi(x_j)$, usiamo una **funzione kernel** che calcola direttamente il prodotto scalare nello spazio trasformato:

$$K(x_i, x_j) = \phi(x_i)^T \phi(x_j)$$

Senza dover esplicitamente calcolare ϕ .

3.3.4.2 Kernel Comuni

1. Kernel Lineare

$$K(x_i, x_j) = x_i^T x_j$$

Nessuna trasformazione; ritorna alla SVM lineare.

2. Kernel Polinomiale

$$K(x_i, x_j) = (x_i^T x_j + 1)^d$$

Trasforma i dati in uno spazio di polinomi di grado d . Utile per pattern polinomiali.

3. Kernel RBF (Radial Basis Function - Gaussian)

$$K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$$

Trasforma in uno spazio di dimensionalità infinita. È il kernel **più comunemente usato** perché:

- Molto flessibile, adatto a pattern complessi
- Un solo iperparametro γ (gamma) da tuning
- Buon compromesso tra complessità e performance

Interpretazione di γ :

- **γ grande:** il kernel focalizza su punti molto vicini (può causare overfitting)
- **γ piccolo:** il kernel considera punti lontani (più generale, rischio underfitting)

4. Kernel Sigmoid

$$K(x_i, x_j) = \tanh(\alpha x_i^T x_j + c)$$

Simile al kernel di una rete neurale. Meno comune.

3.3.4.3 Rappresentazione della Soluzione

La soluzione di SVM è rappresentata come una **combinazione lineare di prodotti kernel**:

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i y_i K(x, x_i)$$

Dove:

- α_i sono i **coefficienti duali** (non direttamente i β_j)
- Solo una **frazione dei punti di training** ha $\alpha_i > 0$ — questi sono i **Support Vectors**

I support vectors sono i punti più informativi per la classificazione; i punti "facili" (ben separati) hanno $\alpha_i = 0$ e vengono ignorati.

3.4 Complessità Computazionale

3.4.1 Training

La complessità di training di SVM dipende dal metodo di ottimizzazione usato:

3.4.1.1 Hard-Margin SVM (Lineare)

$$O(n^2p) \text{ a } O(n^3p)$$

- Metodo **Quadratic Programming (QP)**: tipicamente $O(n^2p)$ o $O(n^3)$ per il solver
- Con n = numero di campioni, p = numero di feature
- La matrice del kernel è $n \times n$, quindi operazioni quadratiche su essa

3.4.1.2 Soft-Margin SVM (con variabili di slack)

$$O(n^2p) \text{ a } O(n^3p)$$

Stessa complessità; la penalità C non cambia la complessità asintotica.

3.4.1.3 Kernel SVM

$$O(n^2p) \text{ a } O(n^3)$$

Per kernel RBF, non è necessario calcolare esplicitamente ϕ (dimensioni potenzialmente infinite):

- Calcolo della matrice kernel: $O(n^2p)$ (comparare ogni coppia di punti)
- Solver QP: $O(n^3)$
- **Total:** $O(n^3)$ (dominato dal solver)

3.4.1.4 Ottimizzazioni Pratiche

Per dataset **grandi** ($n \gg 10.000$), i solver standard diventano intractabili. Soluzioni:

1. **Stochastic Gradient Descent (SGD) SVM:** $O(n p)$ ma con più iterazioni
2. **Online SVM:** aggiorna in streaming
3. **Approximate solutions:** vincolamenti per ridurre n
4. **Distributed training:** parallelizzazione su cluster

3.4.2 Inference (Previsione)

$$O(m \cdot p)$$

Dove m è il **numero di support vectors** ($m \leq n$).

La previsione richiede:

$$f(x_{\text{new}}) = \beta_0 + \sum_{i \in SV} \alpha_i y_i K(x_{\text{new}}, x_i)$$

Se il modello ha molti support vectors ($m \approx n$), inference diventa lenta. Se $m \ll n$ (case ideale), inference è veloce.

3.4.3 Memoria

$$O(n^2 + m \cdot p)$$

- Matrice kernel: $O(n^2)$
- Coefficienti: $O(m)$ (support vectors) + $O(p)$ (per ogni kernel non lineare)

Per dataset molto grandi, la memoria della matrice kernel può essere **proibitiva**.

3.4.4 Scalabilità: Conclusioni

- **Training:** cubic in n — **critico per dataset grandi**
 - Per $n > 100.000$, impraticabile senza approssimazioni
 - Per $n < 10.000$, ragionevole
- **Inference:** dipende da m (numero support vectors)
 - Se m è piccolo: veloce
 - Se $m \approx n$: lento
- **Memoria:** proibitiva per dataset enormi ($n > 1M$)

3.5 Rappresentazione Interna

3.5.1 Struttura del Modello

La rappresentazione interna di SVM è:

Support Vectors: $[x_1, x_2, \dots, x_m]$ (sottinsieme di n punti di training)
Coefficienti: $[\alpha_1, \alpha_2, \dots, \alpha_m]$ e β_0
Kernel: $K(\cdot, \cdot)$ (funzione)

Previsione: $f(x) = \beta_0 + \sum_i \alpha_i y_i K(x, x_i)$

3.5.2 Implicazioni per la Spiegabilità

Contro:

- La previsione dipende da un **subset arbitrario** di punti di training (support vectors)

- Non è immediatamente chiaro **perché** un particolare punto è un support vector
- Per kernel non lineari (es. RBF), la trasformazione ϕ è implicita e non visualizzabile (dimensioni infinite)
- L'interpretazione di α_i non è intuitiva: non corrisponde direttamente a "importanza" della feature

Pro:

- Almeno i **support vectors sono identità**: puoi elencare quali punti di training hanno influenzato la previsione
- Questo è meglio di una rete neurale dove la decisione è completamente dispersa nei parametri
- Puoi analizzare i support vectors localmente per capire il ragionamento

Contrasto con altri modelli:

- **LR/LogR**: coefficienti sono interpretabili per feature (spiegabilità per feature)
- **DT**: cammino è tracciabile (spiegabilità per logica)
- **SVM**: support vectors sono tracciabili (spiegabilità per precedenti)

3.6 Vincoli sui Dati

3.6.1 Linearità (Hard-Margin SVM)

Hard-margin SVM **richiede separabilità lineare esatta**. Se i dati non sono linearmente separabili, il problema di ottimizzazione **non ha soluzione fattibile**.

Soluzione: Soft-margin SVM, che tollera violazioni.

3.6.2 Linearità (Soft-Margin SVM)

Soft-margin SVM è ancora lineare nello spazio dei dati (senza kernel). Se i dati hanno pattern non lineari, si usa kernel SVM.

3.6.3 Scalabilità

Come discusso, **training è $O(n^3)$** , quindi SVM non scala bene a dataset enormi senza approssimazioni.

3.6.4 Bilanciamento delle Classi

SVM può essere influenzato da **classi sbilanciate**:

- L'algoritmo massimizza il margine complessivo, il che potrebbe favorire la classe maggioritaria
- Se una classe è il 95% e l'altra il 5%, il margine "massimo" potrebbe consistere nel classificare tutto come la classe maggioritaria

Soluzioni:

- **Pesi di classe**: assegnare peso inverso alla frequenza (classe rara = peso alto)

- **Tuning di C:** aumentare C per la classe minoritaria
- **Ricampionamento:** oversampling della classe rara, undersampling della maggioritaria

3.6.5 Normalizzazione delle Feature

SVM usa la **distanza euclidea** ($\|x_i - x_j\|$) nel kernel, quindi è **sensibile alla scala**:

- Se una feature va da 0-1 e un'altra da 0-1.000.000, quest'ultima domina
- **Pratica standard:** normalizzare tutte le feature (es. standardizzazione z-score o min-max scaling)

3.6.6 Nessun'altra Assunzione Strutturale

A differenza di LR, SVM **non ha assunzioni** su:

- Normalità delle distribuzioni
- Omoschedasticità
- Relazioni lineari (soprattutto con kernel)

3.7 Capacità Predittive

3.7.1 Punti di Forza

1. Eccellente per Pattern Non Lineari (con Kernel)

- Kernel RBF può approssimare funzioni arbitrariamente complesse
- Spesso outperform modelli lineari su dati reali

2. Robusto a Outlier

- Soft-margin SVM usa slack variables; outlier hanno impatto limitato
- Il margine è costruito basandosi su support vectors, non su tutti i punti

3. Funziona Bene in Alte Dimensioni

- Efficace anche quando $p > n$ (numero di feature > numero di campioni)
- Modelli lineari tendono a overfitting in questi scenari

4. Principi Teorici Solidi

- Massimizzazione del margine ha fondamenti teorici (theory of VC dimension, generalization bounds)
- Garantisce che la soluzione è **globalmente ottimale** (problema convesso)

5. Versatile

- Funziona sia per classificazione che per regressione (SVR)
- Multi-classe mediante one-vs-rest o one-vs-one

3.7.2 Punti di Debolezza

1. Scarsa Interpretabilità

- Specialmente con kernel non lineari, difficile spiegare il ragionamento

2. Sensibilità a Iperparametri

- Scelta di C (soft-margin) e γ (RBF kernel) critica
- Tuning male = performance disastrose

3. Inefficienza su Dataset Grandi

- Complessità training $O(n^3)$ — impraticabile per $n \gg 100.000$
- Anche inference può essere lenta se molti support vectors

4. Non Probabilistico

- Output è una classificazione hard (classe -1 o $+1$)
- Per ottenere probabilità, è necessario post-processing (Platt scaling)

5. Multicollinearità Non Gestita

- Se feature sono fortemente correlate, performance può deteriorarsi
- Richiede feature selection o regolarizzazione

3.8 Metriche per la Confidenza

3.8.1 Metriche di Classificazione Standard

Stesse metriche di LogR e DT:

- **Confusion Matrix:** TP, TN, FP, FN
- **Accuracy:** $(TP + TN) / (TP + TN + FP + FN)$
- **Precision:** $TP / (TP + FP)$
- **Recall/Sensitivity:** $TP / (TP + FN)$
- **Specificity:** $TN / (TN + FP)$
- **F1-Score:** media armonica di Precision e Recall
- **ROC Curve e AUC:** trade-off tra TPR e FPR

3.8.2 Distanza dall'Iperpiano

Misura della **confidenza specifica di SVM**:

$$d_i = y_i \left(\beta_0 + \sum_{j=1}^p \beta_j x_{ij} \right)$$

Questa è la **distanza (segnata) del punto x_i dall'iperpiano**, normalizzata da $||\beta||$:

$$\text{Distanza Normalizzata} = \frac{d_i}{||\beta||}$$

Interpretazione:

- $d_i > 1$: punto ben classificato, lontano dall'iperpiano (confidenza alta)
- $d_i \approx 1$: punto sul margine (confine della decisione)

- $0 < d_i < 1$: punto nella regione di slack (violazione del margine)
- $d_i < 0$: punto classificato erroneamente

Questa distanza è un **proxy per la confidenza della previsione**: punti lontani dall'iperpiano sono predetti con confidenza maggiore.

3.8.3 Platt Scaling (Probabilità Posteriori)

SVM non produce probabilità di default. Per ottenere stime di probabilità $P(y = 1 | x)$, si usa **Platt scaling**:

$$P(y = 1) = \frac{1}{1 + \exp(A \cdot d + B)}$$

Dove A, B sono parametri appresi su un validation set, calibrando la distanza dell'iperpiano a probabilità.

Limite: aggiunge complessità e calibrazione; le probabilità risultanti non hanno la stessa garanzia teorica della regressione logistica.

3.9 Metriche per la Comprensione e Spiegabilità

3.9.1 1. Identificazione dei Support Vectors

Il set dei support vectors è la **feature più interpretabile** di SVM:

Support Vectors: [istanza_3, istanza_15, istanza_42, ...]

Rappresentano i punti di training che hanno **influenzato veramente la decisione**. Punti ben separati hanno $\alpha_i = 0$ e sono ignorati.

Utilizzo pratico:

- Analizzare i support vectors per capire quali punti sono "difficili" (contorno decisionale)
- Visualizzare support vectors vs punti di background per ispezionare il modello
- Se molti support vectors = modello complesso; se pochi = modello semplice

3.9.2 2. Pesi dei Support Vectors (α_i)

I coefficienti $\alpha_i > 0$ associati ai support vectors indicano il loro "peso" nella decisione:

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i y_i K(x, x_i)$$

Interpretazione:

- α_i grande = il support vector x_i ha grande influenza
- α_i piccolo = il support vector x_i ha piccola influenza

Caveat: interpretazione limitata per kernel non lineari, perché non è immediatamente chiaro cosa stia facendo il kernel.

3.9.3 3. Feature Importance (via Permutation o SHAP)

SVM non produce naturalmente feature importance. Soluzioni:

Permutation Importance:

- Permutare casualmente una feature nei dati di test
- Misurare il degrado in performance
- Grandi degni = feature importante

SHAP (SHapley Additive exPlanations):

- Assegnare ogni output tra le feature usando valori Shapley
- Applicabile a qualsiasi modello (SVM, NN, etc.)
- Computazionalmente costoso pero produce spiegazioni teoricamente fondate

3.9.4 4. Visualizzazione della Separazione (2D/3D)

Per dataset a bassa dimensionalità, visualizzare:

- Punti di training e loro classi
- Iperpiano di decisione (retta in 2D, piano in 3D)
- Margine (linee parallele a distanza $M = 1 / ||\beta||$)
- Support vectors (segnati diversamente)

Limite: per $p > 3$, non è possibile visualizzare direttamente. Usare proiezione PCA per ridurre dimensioni e visualizzare.

3.9.5 5. Analisi Locale: Contrastive Explanation

Per una previsione di istanza x_{new} :

- Trovare i support vectors più vicini (usando la metrica del kernel)
- Mostrare come differiscono da x_{new}

Esempio:

Istanza X è classificata come Positiva perché simile al Support Vector SV_5 (che è Positivo), diversa dal Support Vector SV_12 (che è Negativo).

Differenze principali:

- X.Income > SV_5.Income (simile)
- X.Age < SV_12.Age (dissimile)

3.10 Limiti di Predizione

3.10.1 Non Probabilistico per Default

SVM output una classificazione hard (+1 o -1), non una probabilità:

- Non ottenere stime naturali di incertezza
- Platt scaling produce probabilità, ma richiede calibrazione aggiuntiva

Conseguenza: per problemi dove la probabilità importa (es. medicina), SVM deve essere post-processato.

3.10.2 Sensibilità al Tuning di Iperparametri

Performance di SVM è **molto sensibile** a:

- **C (soft-margin):** piccolo C = underfitting; grande C = overfitting
- **γ (RBF kernel):** piccolo γ = generale; grande γ = specifico (overfitting)

Conseguenza: il tuning di iperparametri (grid search, random search, Bayesian optimization) è **mandatorio**, il che aggiunge complessità.

3.10.3 Sensibilità alla Normalizzazione delle Feature

Senza normalizzazione, feature con scale grandi dominano. SVM è **molto sensibile** a questo.

Mitigazione: normalizzare sempre le feature (standardizzazione z-score è standard).

3.10.4 Performance su Dataset Sbilanciate

Come discusso, SVM tende a favorire la classe maggioritaria.

Mitigazione: usare class weights (inversamente proporzionali alla frequenza).

3.10.5 Inefficienza su Dataset Grandi

Complessità $O(n^3)$ rende SVM **impraticabile per $n \gg 100.000$** .

Mitigazioni possibili:

- Usare approssimazioni (Nyström approximation, stochastic gradient descent)
- Sottocampionare per training, validare su full dataset
- Usare alternative scalabili (logistic regression, gradient boosting)

3.11 Limiti di Spiegabilità

3.11.1 Opacità della Trasformazione Non Lineare (Kernel)

Con kernel RBF, i dati sono trasformati in uno **spazio di dimensionalità infinita**:

$$K(x, x') = \exp(-\gamma ||x - x'||^2)$$

Questa trasformazione $\phi(x)$ è **implicita e non visualizzabile**. Conseguenze:

- Non puoi interpretare cosa stia facendo il modello nello spazio trasformato
- L'effetto delle feature sulla previsione è **nascosto dietro il kernel**
- Spiegare una previsione richiede capire il kernel, che è non intuitivo per i non-esperti

3.11.2 Interpretazione Limitata di α_i

I pesi α_i dei support vectors non hanno interpretazione diretta:

- Non significano "questa feature è importante"

- Non significano "questo support vector cambia la probabilità del 10%"
- Sono artefatti dell'ottimizzazione quadratica, non quantità interpretabili

3.11.3 Dipendenza da Support Vectors Arbitrari

Quali punti diventano support vectors dipende dalla **geometria dello spazio** e dall'**iperparametro C**:

- Cambiar $C \rightarrow$ diversi support vectors
- Cambiar dataset leggermente $\rightarrow$ diversi support vectors
- Non è evidente **perché** un punto è un support vector (è dovuto alla posizione geometrica, ma questa non è facilmente comunicabile)

3.11.4 Scarsa Tracciabilità Globale

Mentre puoi identificare i support vectors, **tracciare il ragionamento globale è difficile**:

- Non puoi dire "la feature X è il driver principale della classificazione" (come in LR)
- Non puoi tracciare una sequenza logica (come in DT)
- Rimane una "scatola nera" anche dopo aver visto i support vectors

3.11.5 Nessuna Spiegazione Naturale per Feature Importanza

A differenza di:

- **LR**: coefficienti diretti
- **DT**: feature importance dai split

SVM non produce feature importance naturale. Devi usare tecniche esterne (SHAP, permutation importance).

3.12 Confronto con altri Algoritmi

3.12.1 Vs. Logistic Regression

Aspetto	LogR	SVM
Output	Probabilità	Classe (distanza)
Interpretabilità	★★★★★ Alta	★★★ Bassa
Pattern Non-Lineare	★ (no kernel)	★★★★★★ (con kernel RBF)
Scalabilità	★★★★★ $O(n \cdot p \cdot k)$	★★★ $O(n^3)$
Teoria	Probabilistica	Geometrica (margine)
Quando usare	Interpretabilità critica	Pattern non-lineare complesso

3.12.2 Vs. Decision Tree

Aspetto	DT	SVM
Interpretabilità	★★★★★★ Altissima	★★★ Bassa
Pattern Non-Lineare	★★★★★★ Naturale	★★★★★★ Con kernel
Scalabilità	★★★★★ $O(n \cdot p \cdot \log n)$	★★★ $O(n^3)$
Robustezza Outlier	★★★ (sensibile)	★★★★★★ (robusto)
Stabilità	★★★ (instabile)	★★★★★★ (stabile)
Quando usare	Massima interpretabilità	Pattern complessi + stabilità

3.12.3 Vs. Neural Networks

Aspetto	NN	SVM
Capacità	★★★★★ Illimitata	★★★★★ Molto buona
Interpretabilità	★ Scatola nera	★★ Leggermente migliore
Scalabilità	★★★★★ Eccellente	★★ Scarsa
Data Requirement	★ (molto dati)	★★★★★ (meno dati)
Training	Lungo, GPU-friendly	Veloce ma $O(n^3)$
Teoria	Inesistente	Solida
Quando usare	Big data + capacità illimitata	Data medio + interpretabilità

3.12.4 Vs. Random Forest / Gradient Boosting

Aspetto	Ensemble	SVM
Performance	★★★★★ Spesso migliore	★★★★★ Buona
Interpretabilità	★★ (migliore di NN)	★★ Simile
Scalabilità	★★★★★ (meglio di SVM)	★★
Robustezza Iperparametri	★★★★★ Meno sensibile	★★ Molto sensibile
Implementazione	Easy (sklearn)	Easy (sklearn)
Quando usare	Predizione è priorità #1	Stabilità, spiegabilità, dati piccoli

3.13 Varianti e Estensioni

3.13.1 1. Multi-Class SVM

SVM è originariamente binario. Per multi-classe:

One-vs-Rest (OvR):

- Creare K classificatori binari ($K = \text{numero di classi}$)
- Ogni classificatore separa una classe dal resto
- Classe finale = massima confidenza

One-vs-One (OvO):

- Creare $K(K-1)/2$ classificatori binari (uno per coppia di classi)
- Voting: conta quante volte ogni classe è predetta
- Più computazionalmente costoso, spesso più accurato

3.13.2 2. Support Vector Regression (SVR)

Estensione di SVM per **regressione** (output continuo):

- Anziché maximizzare il margine di classificazione, minimizzare l'errore di predizione
- Introduce un **'epsilon-insensitive loss'**: errori piccoli $< \epsilon$ sono tollerati
- Utile per predizione continua con gestione robusta di outlier

3.13.3 3. Uno-Classe SVM (One-Class SVM)

Per problemi di **anomaly detection** / **outlier detection**:

- Imparare un "confine" intorno ai dati normali
- Punti fuori da questo confine sono anomalie
- Usa soft-margin con una frazione ν di campioni come support vectors

3.13.4 4. Relevance Vector Machine (RVM)

Variante Bayesiana di SVM:

- Produce probabilità di output come LogR
- Spesso meno support vectors (più sparso)
- Trade-off: meno efficienza di SVM, più interpretabilità

3.14 Raccomandazioni Pratiche

3.14.1 Quando Usare SVM

Ottime scelte:

- Dataset di piccolo-medio volume ($n < 100.000$)
- Pattern non lineari complessi
- Necessità di stabilità e robustezza

- Dati ad alta dimensionalità ($p > n$)
- Outlier presenti nel dataset

✗ Cattive scelte:

- Dataset molto grandi ($n > 1.000.000$)
- Interpretabilità è critica per decisioni (es. medicina, finanza regolata)
- Tuning di iperparametri non è possibile
- Output probabilistico naturale necessario

3.14.2 Tuning Consigliate

1. **Normalizzare sempre le feature** (z-score o min-max)
2. **Griglia di ricerca per C e γ :**
 - C : [0.1, 1, 10, 100, 1000]
 - γ : [0.001, 0.01, 0.1, 1, 'auto']
3. **Cross-validation:** 5-10 fold per valutare
4. **Class weights:** se classi sbilanciate, usare pesi inversi
5. **Kernel:** iniziare con RBF; provare polinomiale se pattern sospetto

3.15 Prompt

3.16 ...

Chapter 4

Design and code implementation

In this chapter we will describe the design and the implementation of the system, starting from the requirements and the technologies used, to the design and coding of the system.

4.1 Requirements

Before implementing the system, it is crucial to define what the goals and the requirements of the project are. For a more precise description this requirements are categorized in functional (F), qualitative (Q) and constraint (V) requirements, and each of them is classified as necessary (N), desirable (D) or optional (O).

Requirement	Description
RFN-1	The system must allow the user to choose the dataset to analyse
RFN-2	The system must allow the user to choose the algorithm for the analysis
RFN-3	The system must allow the user to choose the level of detail for the analysis (example SHAP/No-SHAP)
RFN-4	The system must allow the user to choose whether to execute the LLM analysis or not
RFN-5	The system must allow the user to access the result of the analysis in a clear and organized way
RFN-6	The system must allow the user to know the reason in case of an error during the execution of the analysis
RFO-1	The system should be accessible via graphical user interface (No-CLI)
RQN-1	The system should be open to the use of different datasets, algorithms

Requirement	Description
RQN-2	The system should be open to the use of different explainability and analysis techniques
RQN-3	The system should be open to the use of different LLMs for the analysis
RQD-1	The system should process the data and execute the analysis in a reasonable time frame, including the LLM analysis the pipeline should not take more than 15 minutes to execute

Table of requirements for the analysis pipeline.

4.2 Technologies and tools

Considering the requirements defined in the previous section and the context of the project, the following technologies and tools were chosen for the implementation of the system:

4.2.1 Python

Python has been chosen as the main programming language for the implementation of the system due to its versatility, ease of use, and the wide range of libraries and frameworks available for data analysis, machine learning, and natural language processing.

4.2.2 Markdown

Markdown has been chosen for the collection of the LLM responses to generate the final report. It was chosen for its simplicity and readability, as well as its compatibility with various tools and platforms for documentation and reporting.

4.2.3 Pandas, Matplotlib, Seaborn, Numpy, SHAP

These Python libraries have been chosen for the vast range of functionalities. Pandas for data manipulation and analysis, Matplotlib and Seaborn for data visualization, Numpy for numerical computations, and SHAP for explainability analysis ([Section 2.2.3](#)).

4.2.4 Scikit-learn

Scikit-learn provides implementations of most of the algorithms in exam, without the need to implement them from scratch, allowing to focus on the analysis and explainability aspects of the project.

Moreover, it provides a variety of metrics for evaluating the performance of the models, which is crucial for the analysis.

The models that were not offered by scikit-learn were implemented using the original implementations provided by the authors of the algorithms, for example for XGBoost[8].

4.2.5 Optuna

Since a lot of algorithms benefits from hyperparameter tuning, Optuna was chosen for its efficiency and ease of use in automatically optimizing the hyperparameters of machine learning models.

4.2.6 Streamlit

Streamlit was chosen for the implementation of the **Graphical User Interface (GUI)** of the system. It allows to quickly and easily create interactive web applications for data analysis and machine learning, which is essential for providing a user-friendly interface for the system.

4.2.7 GitHub

GitHub was chosen for the version control of the project, for an effective codebase management.

4.2.8 Typst

Typst was chosen as principal tool for the notes and documentation. Thanks to its flexibility, modern design, powerful features, effective rendering of the PDFs, Typst

4.3 Design

The design of the entire pipeline was made with the goal of an extendable pipeline, providing the developer the possibility to easily add new algorithms, datasets, analysis techniques and **LLMs**. For this reason, the design of the system is modular, with each component of the pipeline being independent and interchangeable.

The main components of the system are:

IMMAGINE DA INSERIRE CON STARUML

1. **Orchestrator**: coordinates the pipeline, invoking the different components in the correct order and passing the necessary data between them.
2. **Selector**: allows the user to choose the dataset, the algorithm, the level of detail for the analysis and whether to execute the **LLM** analysis or not.
3. **Algorithm**: implements the machine learning algorithm and provides the necessary functionalities for training, prediction and evaluation.
4. **LLM Request Manager**: manages the requests to the **LLM**, including the generation of prompts and the collection of responses.

4.3.1 Guiding principles

For the best possible design of the system have been used the principles of object-oriented programming and some design patterns. First of all the SOLID principles [9], [10]. In particular:

1. **Single Responsibility Principle:** each class has a single responsibility, making the code more modular and easier to maintain. A clear example is the modular structure cited just before.
2. **Open/Closed Principle:** each implementation of the abstract `baseMLAlgo` can expand the possibility of the base class, adding new functionality without modifying the existing code.
3. **Liskov Substitution Principle:** the implementation of the `baseMLAlgo` substitute the abstract class, applying their version of the functions and modifying the analysis.
4. **Interface Segregation Principle:** the interfaces of the different components are designed to be specific to their functionality, avoiding unnecessary dependencies between them. For example, the `LLMRequestManager` has a specific interface for managing the requests to the LLM, without depending on the implementation of the algorithms or the orchestrator.
5. **Dependency Inversion Principle:** the high-level modules (`orchestrator`) do not depend on low-level modules (algorithms, LLM request manager), but both depend on abstractions. For example, the orchestrator depends on the abstract `baseMLAlgo` and `LLMRequestManager`, allowing to easily switch between different implementations of the algorithms and the LLM request manager without modifying the orchestrator.

4.3.2 Applied design patterns

The design of the system also incorporates some design patterns to solve common problems and improve the structure of the code. In particular, the following design patterns were applied:

4.3.2.1 Strategy

The strategy pattern was applied to the implementation of the algorithms, allowing to easily switch between different algorithms without modifying the code of the `orchestrator`. Each algorithm implements the same interface defined by the abstract `baseMLAlgo`, allowing to use them interchangeably in the orchestrator.

IMMAGINE DA INSERIRE CON STARUML

4.3.2.2 Template method

The base class `baseMLAlgo` defines the template method `run_analysis()`, which outlines the steps of the analysis process, while the concrete implementations of the algorithms provide their specific implementation of each step. This allows to have a consistent structure for the analysis while allowing for flexibility in the implementation of each algorithm.

IMMAGINE O SNIPPET DI CONFRONTO TRA BASEMLALGO E UNA IMPLEMENTAZIONE CONCRETA

4.3.2.3 Facade

The facade pattern was applied to the *orchestrator*, providing a single entrance *run_pipeline()* for managing the entire process of generating prompts, sending requests to the LLM and collecting responses, hiding the complexity of the underlying implementation from the rest of the system.

4.3.2.4 Registry

The registry pattern was applied to the management of the algorithms, allowing to easily register new algorithms and access them through a central registry. This allows to easily expand the system with new algorithms without modifying the existing code.

4.4 Code implementation

Chapter 5

Prompt Engineering

Breve introduzione al capitolo

Chapter 6

Conclusion

6.1 Consuntivo finale

6.2 Raggiungimento degli obiettivi

6.3 Conoscenze acquisite

6.4 Valutazione personale

Bibliography

- [1] Christoph Molnar, *Interpretable Machine Learning: A Guide for Making Black Box Models Explainable*. Leanpub, 2025.
- [2] Zizhao Hu, Mohammad Rostami, and Jesse Thomason, “Expert Personas Improve LLM Alignment but Damage Accuracy: Bootstrapping Intent-Based Persona Routing with PRISM,” 2026.
- [3] “Which Table Format Do LLMs Understand Best?.” [Online]. Available: <https://www.improvingagents.com/blog/best-input-data-format-for-llms/>
- [4] Jason Wei *et al.*, “Chain of Thought Prompting Elicits Reasoning in Large Language Models,” 2023.
- [5] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa, “Large Language Models are Zero-Shot Reasoners,” 2023.
- [6] Taowen Wang *et al.*, “M²PT: Multimodal Prompt Tuning for Zero-Shot Instruction Learning,” 2024.
- [7] “P-Value: What It Is, How to Calculate It, and Examples.” [Online]. Available: <https://www.investopedia.com/terms/p/p-value.asp>
- [8] Tianqi Chen and Carlos Guestrin, “XGBoost: A Scalable Tree Boosting System,” *KDD '16*, 2016.
- [9] Robert C. Martin, “Design principles and design patterns,” 2002. [Online]. Available: https://objectmentor.com/resources/articles/Principles_and_Patterns.pdf
- [10] Robert C. Martin, *Clean Code: A Handbook of Agile Software Craftsmanship*. Pearson, 2008.
- [11] Zhen "Leo" Liu, *Artificial Intelligence for Engineers: Basics and Implementations*. Springer, 2024.
- [12] Candice Bentéjac, Anna Csörgő, and Gonzalo Martínez-Muñoz, “A Comparative Analysis of XGBoost,” *Artificial Intelligence Review*, vol. 54, no. 3, 2019.
- [13] Christopher M. Bishop, *Pattern Recognition and Machine Learning*. Springer, 2006.
- [14] G. James, D. Witten, T. Hastie, R. Tibshirani, and J. Taylor, *An Introduction to Statistical Learning: with Applications in Python*. Springer, 2023.
- [15] V. K. Xidong Wu, *Top 10 Algorithms in Data Mining*. CRC Press, 2009.
- [16] Lior Rokach and Oded Maimon, *Data Mining with Decision Trees: Theory and Applications*. World Scientific Publishing, 2008.
- [17] “SHAP Documentation.” [Online]. Available: <https://shap.readthedocs.io/en/latest/>

- [18] Xavier Amatriain, “Prompt Design and Engineering: Introduction and Advanced Methods,” 2024.

Chapter 7

Glossary

AI – Artificial Intelligence: Field of computer science that focuses on creating systems capable of performing tasks that typically require human intelligence, such as learning, reasoning, problem-solving, and understanding natural language. [1](#)

categorical features – Categorical Features: Features in a dataset that represent discrete categories or groups, rather than continuous numerical values, and often require special handling in machine learning algorithms. [4](#), [13](#), [15](#), [20](#)

classification – Classification: A type of machine learning problem in which the goal is to predict discrete class labels based on input features. [i](#), [2](#)

computational complexity – Computational complexity: A measure of the amount of computational resources (time and space) that an algorithm requires to run, often expressed in terms of the size of the input data. [4](#), [12](#), [13](#), [13](#), [13](#), [13](#)

correlation matrix – Correlation Matrix: A table showing the correlation coefficients between pairs of features in a dataset, used to identify relationships and potential multicollinearity. [14](#), [21](#)

XAI – Explainable Artificial Intelligence: A field of artificial intelligence that focuses on creating models and systems that can provide clear and understandable explanations for their decisions and actions. [i](#), [i](#), [2](#), [9](#)

FN – False Negatives: Instances incorrectly predicted as negative. [22](#), [22](#)

FP – False Positives: Instances incorrectly predicted as positive. [22](#), [22](#)

GD – Gradient Descent: An optimization algorithm used to minimize the objective function in machine learning models by iteratively adjusting the model's parameters in the direction of the steepest descent. [12](#), [13](#), [13](#), [14](#), [19](#), [19](#)

GUI – Graphical User Interface: A user interface that allows users to interact with a computer system through graphical elements such as windows, icons, and buttons, rather than text-based commands. [44](#)

LLM – Large Language Model: A type of artificial intelligence model that is trained on vast amounts of text data to understand and generate human-like language. [i](#), [i](#), [2](#), [7](#), [7](#), [7](#), [9](#), [12](#), [44](#), [44](#), [44](#), [45](#), [46](#)

MAE – Mean Absolute Error: A common metric for evaluating the performance of regression models, calculated as the average of the absolute differences between predicted and actual values. [iv](#), [16](#)

margin – Margin: The distance between the decision boundary (hyperplane) and the closest data points in a Support Vector Machine, which the algorithm seeks to maximize for better generalization. [25](#), [25](#), [26](#)

ML – Machine Learning: A subset of artificial intelligence that involves training algorithms to learn patterns from data and make predictions or decisions without being explicitly programmed. [i](#), [1](#)

objective_function – Objective Function: A mathematical function that an algorithm optimizes during training, which defines the goal of the learning process and guides the model's adjustments to minimize error or maximize performance. [4](#)

outlier: A data point that deviates significantly from the majority of the data, which can have a disproportionate influence on machine learning models and may require special handling. [15](#)

PE – Prompt Engineering: The process of designing and optimizing prompts to effectively communicate with language models and elicit desired responses. [2](#)

regression – Regression: A type of machine learning problem in which the goal is to predict continuous numerical values based on input features. [i](#), [2](#)

RMSE – Root Mean Squared Error: A common metric for evaluating the performance of regression models, calculated as the square root of the average of the squared differences between predicted and actual values. [iv](#), [16](#)

SHAP – SHapley Additive exPlanations: A method for explaining the output of machine learning models by attributing the prediction to each feature. [i](#)

SGD – Stochastic Gradient Descent: An optimization algorithm that updates the model's parameters using a randomly selected subset of the training data, which can lead to faster convergence and better generalization. [13](#), [20](#)

TN – True Negatives: Instances correctly predicted as negative. [22](#), [22](#)

TP – True Positives: Instances correctly predicted as positive. [22](#), [22](#)